

# **Compliance Aware, Agentic Pitch Book Generation For Bankers**

A template-aware PowerPoint generation system with agentic data retrieval for investment banking.

A Queen Mary Final Year Dissertation

**Andrei Rizea**

Student ID: 220300153

Supervised by: Manesha Peiris

Programme of Study:

Digital & Technology Solutions (Software Engineering)

Module: IOT635W

Queen Mary University of London

January 21, 2026

## **Abstract**

# Contents

<b>Abstract</b>	<b>1</b>
<b>Introduction, Scope and Context</b>	<b>5</b>
Where Analyst Time Goes .....	5
Problem Analysis .....	6
Gap Analysis .....	7
Aims and Objectives .....	8
Scope and Success Criteria .....	9
<b>Methodology and Project Plan</b>	<b>10</b>
Development Methodology .....	10
Requirements Analysis .....	11
Business Requirements .....	11
Functional Requirements .....	12
Non-Functional Requirements .....	12
Technology Stack .....	13
Programming Language and Framework .....	13
AI and LLM Integration .....	14
Data Sources and APIs .....	15
Frontend and Deployment .....	16
Reproducibility .....	17
Evaluation Methodology .....	17
Project Plan .....	18
Risk Assessment and Mitigation .....	20
Ethics, Data Governance and Compliance .....	21
Project Tracking .....	22
Reusability .....	22
<b>Preliminary Research and Design Documentation</b>	<b>23</b>
<b>Project Implementation and Outcomes</b>	<b>26</b>
<b>Evaluation and Conclusions</b>	<b>27</b>
<b>References</b>	<b>28</b>
<b>Appendices</b>	<b>29</b>

## List of Figures

1	Project Gantt Chart (19 January to 20 April 2026)	19
---	---	----

# List of Tables

1	Comparison of Development Methodologies . . . . .	10
2	Functional Requirements Specification . . . . .	12
3	Non-Functional Requirements Specification . . . . .	13
4	Programming Language Comparison . . . . .	14
5	Large Language Model Comparison . . . . .	15
6	External Data Sources . . . . .	16
7	Frontend Framework Comparison . . . . .	16
8	Deployment Platform Comparison . . . . .	17
9	Project Milestones and Deliverables . . . . .	20
10	Risk Assessment and Mitigation Strategies . . . . .	21

# Introduction, Scope and Context

Knowledge workers are people whose jobs involve handling information rather than physical goods: analysts, consultants, lawyers, accountants, and similar professionals. McKinsey found they spend a fifth of their time just searching for information (McKinsey Global Institute 2012). Investment banking (IB) is worse. By some estimates, document preparation eats up half to three-quarters of a junior banker's week, and pitch books (PBs) account for a big chunk of that (Corporate Finance Institute 2025; Wall Street Oasis 2024). When Goldman Sachs surveyed its first-year analysts in 2021, the average was 98 hours a week, and preparing documents was a major driver of the burnout they reported (BBC News 2021). Yet most PBs look nearly identical to one another. Change the company name and update the numbers, and you could reuse last month's deck. This repetitiveness makes them good candidates for automation.

IB firms advise companies on mergers, acquisitions, capital raising, and other financial transactions. The work is high-stakes. A single deal might be worth billions of dollars, and the documents that support these deals need to be accurate, professional, and persuasive. PBs are the primary vehicle for communicating analysis to clients and internal stakeholders. They summarise a company's financial position, compare it to peers, and make the case for why a particular transaction makes sense. Getting these documents right matters, but getting them done quickly matters too.

The industry has tried various approaches to this problem. Some firms have built internal document libraries where analysts can search for past work. Others have hired dedicated presentation specialists to handle formatting. A few have experimented with automation tools. None of these solutions have fully addressed the underlying issue: too much skilled human time goes toward work that does not require skill.

## Where Analyst Time Goes

IB sells expertise rather than physical goods, so you would expect analyst time to go toward analysis and client relationships. Often it does not.

A typical PB contains an executive summary, company overview, market analysis, valuation work, transaction comparables, and next steps. These sections appear in nearly every deck, along with the same content types: company descriptions, financial tables, market charts, and deal diagrams. The data inside changes, but the shell around it rarely does.

Analyst tasks fall into two buckets. The first holds work that genuinely needs a trained banker: financial modelling, valuation analysis, due diligence, and client advisory. The second holds everything else: hunting down data, reformatting slides, and fixing fonts. Too many hours go into the second bucket.

The process typically works like this. A managing director wins a new engagement and assigns it to an analyst team. The analysts start by digging through the firm's document archives, looking for similar deals they can use as templates. They pull financial data from Bloomberg terminals, public filings, and

company reports. They copy and paste this information into PowerPoint, reformatting it to match the firm's visual standards. Then they send a draft up the chain for review. Senior bankers mark it up with corrections, and the analysts revise and resubmit. This cycle might repeat several times before the document is ready for the client.

Each step in this process has friction. The document archive search is often slow and incomplete. Data entry is error-prone. Formatting is tedious. Review cycles eat up time on both ends. The result is that a significant portion of analyst hours goes toward activities that could, in theory, be automated.

Banks pay junior analysts well, so time spent on low-skill tasks represents a poor return on that investment. Shifting even some of that document preparation time back toward analytical work would improve the economics. This is not about replacing analysts. It is about freeing them to do analyst work.

Research supports this view. Radhakrishna et al. found that well-designed knowledge management systems boosted productivity by 20 to 25 percent (Radhakrishna et al. 2024). Document production was identified as a prime candidate for automation, provided humans retained control over judgement calls.

## Problem Analysis

Markus defined knowledge reuse in 2001, and her framework fits PB production well (Markus 2001). She identified several reuse patterns: drawing on past work, borrowing from colleagues, finding subject matter experts, and mining old documents. PB creation involves all of these, often in the same afternoon.

Three things get in the way.

First, too much time goes to low-value work. Analysts spend hours searching through old decks, extracting useful content, and reformatting slides to fit the current template. None of that requires analytical thinking, yet it crowds out the work that does.

Second, institutional knowledge tends to disappear. Researchers distinguish between tacit knowledge, which stays in people's heads, and explicit knowledge, which gets written down (Dalkir 2011). PBs themselves are explicit since they capture analysis in a shareable format. But the knowledge of how to make a good PB remains tacit: which templates work for which situations, how to structure the narrative, what separates a persuasive pitch from a forgettable one. New analysts pick this up by watching senior colleagues, not by reading a manual.

Third, onboarding takes longer than it should. Each bank has its own templates, preferred data sources, writing conventions, and quality expectations. Junior analysts absorb these through trial and error, which is slow for them and a drain on the senior bankers who review their work. A new analyst might spend weeks learning which template to use for which situation, how to format financial tables correctly, and what level of detail is expected in different sections. This knowledge exists somewhere in the organisation, but finding it and absorbing it takes time.

These three problems reinforce each other. When knowledge is scattered, people waste time searching for it. When standards live in people's heads rather than in documents, output quality varies and review

cycles drag on. When training happens informally, experienced staff lose hours to coaching that could have gone toward billable work. A system that tackled even one of these issues would help. One that addressed all three could change how PB production works entirely.

## Gap Analysis

Automated presentation tools have existed for years. Early versions converted text to slides using layout algorithms, but although the content was organised, the slides looked bad and nobody wanted to use them (Zheng et al. 2025). Template-based systems came next, where you define the structure upfront and the system fills in the blanks. That fixed the visual problems but made everything rigid.

PPTAgent takes a different approach (Zheng et al. 2025). It analyses reference presentations, extracts their structural patterns, and applies those patterns to new content. Because it learns from real examples rather than following hardcoded rules, it beat older text-to-slide systems on both content quality and design quality in testing.

The catch is that PPTAgent targets general-purpose presentations, but IB PBs have specific conventions that generic tools miss. They have financial tables with specific column layouts, transaction comparables formatted in particular ways, and market positioning charts that follow visual conventions the industry recognises. Generic tools miss these conventions entirely, and in IB, getting the template wrong undermines credibility. You could have the best analysis in the world, but if the deck looks off, nobody will take it seriously.

Commercial tools share this flaw. Beautiful.ai, Tome, and Gamma can all generate slides from a prompt, but the output is generic and every slide needs rework to match firm standards. None of them connect to financial data sources either, so analysts still have to gather numbers and enter them by hand. These tools work well for general business presentations, but they were not built for the specific demands of financial services.

The challenge is that IB PBs require both domain knowledge and visual precision. The content needs to reflect financial expertise. The formatting needs to match exacting corporate standards. Generic AI tools typically handle one or the other, but not both. A system combining financial data retrieval, domain-appropriate content, and template-aware assembly would fill this gap.

Here is the gap. Current systems either produce generic output that needs heavy cleanup, or they require large training datasets to learn industry-specific patterns. Most banks lack those datasets. On the data side, no existing tool connects slide generation to financial APIs. Analysts still pull numbers from Yahoo Finance or public filings, copy them into Excel, format them, and paste them into PowerPoint. The data retrieval and the document generation remain separate manual steps. What nobody seems to have built yet is a system that can take a single reference template, extract its styling and layout rules programmatically, pull live financial data from APIs, and populate tables with current numbers while matching the template exactly.

Agentic AI architectures look promising for this (IBM 2024). Instead of one model doing everything

in a single pass, the system breaks tasks into steps and uses the right tool for each one. Wang et al. surveyed LLM-based agents and found that progress has been rapid (Wang et al. 2024). McKinsey estimates generative AI could deliver \$340 billion annually in banking (McKinsey & Company 2024). Document-heavy workflows are a natural fit.

My project addresses this gap. The pipeline extracts layouts, colour palettes, and placeholder positions from a template using python-pptx, uses an LLM to plan content, then assembles slides matching the original style. To coordinate these steps, I am using existing agentic frameworks rather than building custom orchestration. RAG against past PBs is a stretch goal.

The approach differs from prior work in two important ways. First, it prioritises template fidelity over content generation. The system starts with the visual constraints and works backward to content, rather than generating content and hoping the formatting works out. Second, it uses an agentic architecture where specialised components handle distinct parts of the task. One component analyses templates, another retrieves data, another plans content, and another assembles slides. This separation makes the system easier to test and modify.

## Aims and Objectives

The aim is to design, build, and test a proof-of-concept (POC) demonstrating how AI-powered document generation can tackle the knowledge reuse problems described above. The system must respect institutional templates throughout. If it generates slides that do not match the house style, the whole thing fails.

On the business side, I want a workflow where an analyst can provide minimal input and get back a solid first draft. That would cut the time spent on formatting grunt work while keeping humans in control of the actual analysis. The target users are junior analysts, the people who currently spend most of their hours on document preparation.

On the learning side, I want to understand how to connect LLM orchestration with programmatic document generation. That means getting into agentic architectures, learning how to extract patterns from templates, and working out how to combine multiple AI capabilities into something that functions end to end. I am also interested in human-in-the-loop (HITL) design, since in a regulated industry like finance, AI tools need to support professional judgement rather than try to replace it. Banks face heavy regulation, and any AI system deployed there must be transparent, auditable, and controllable by human operators.

Five objectives structure the work:

1. Build a template analysis component that takes a reference PowerPoint file and extracts its layout structures, colour palettes, font specifications, and placeholder positions.
2. Build an agentic data retrieval layer that automatically pulls company financials, market data, and regulatory filings from Yahoo Finance and SEC EDGAR.

3. Build an LLM-powered content planning module that produces structured slide specifications adapted to different PB types while keeping the overall narrative coherent.
4. Build a slide assembly component that takes the content plan and template patterns and produces a PowerPoint file meeting formatting standards.
5. Evaluate the system by measuring generation speed, template matching, and content quality.

## Scope and Success Criteria

The POC takes user input (company, transaction type, dates, reference template, PB type) and produces a formatted PowerPoint. Target PB types are company overviews, market updates, and transaction summaries. Full valuation presentations are out of scope.

Design choices: established agentic frameworks rather than custom architecture, programmatic template extraction with python-pptx rather than vision models, public APIs only, and RAG as a stretch goal. Out of scope: real-time data feeds, production infrastructure, compliance validation. Deal positioning and valuation judgements stay with analysts. Practical limits: one academic term, standard hardware, student API budget.

Success criteria:

- Generate complete PB draft within five minutes
- Match reference template formatting
- Accurate financial data from source APIs
- Content structure following IB conventions, verified against precedent examples

One principle runs through all of this: HMTL design. The system produces drafts for human review, not finished products. The goal is a solid starting point rather than a blank slide. This approach acknowledges that AI systems are not yet capable of the nuanced judgement that financial analysis requires, but they can handle the mechanical work that absorbs too much human attention.

# Methodology and Project Plan

This section covers how I will approach development, what the requirements are, which technologies I have chosen and why, the project schedule, and how I am thinking about risks and ethics. Reproducibility matters throughout. I want anyone reading this work to be able to understand my choices, replicate my setup, and reproduce my results.

## Development Methodology

Software development methodologies range from plan-driven approaches like Waterfall to adaptive approaches like Agile (Sommerville 2016). Choosing the right methodology depends on requirement stability, project complexity, stakeholder availability, and team size. Table 1 compares methodologies against criteria relevant to this project.

Criteria	Waterfall	Scrum	Iterative Prototyping
Requirements stability	Requires fixed requirements upfront	Accommodates changing requirements	Works well when requirements change
Feedback loops	Late feedback after implementation	Sprint reviews every 2-4 weeks	Continuous feedback through prototypes
Risk management	Risks discovered late	Regular risk reassessment	Early risk identification through prototypes
Solo developer suitability	Moderate	Low (designed for teams)	High
Research integration	Poor	Moderate	Excellent

Table 1: Comparison of Development Methodologies

I will use a hybrid methodology combining iterative prototyping with structured research engineering practices. Pure Waterfall cannot accommodate the exploratory nature of this work. Full Scrum adds overhead without benefit for a solo developer. Iterative prototyping lets me adapt as I learn, while structured practices keep the work repeatable.

The choice of iterative prototyping reflects the nature of the problem. I am building a system that combines several technologies in ways that have not been done before. Some technical decisions will only make sense after I have tried them. For example, I cannot know in advance whether a particular prompt structure will produce good slide content until I have tested it with real data. The methodology needs to accommodate this kind of learning.

At the same time, this is academic work that others should be able to evaluate and build upon. That means

I need to document my decisions, track my progress, and make my process transparent. The structured practices make that possible.

Each two-week iteration will produce a working prototype targeting defined requirements. At each cycle's end, I will conduct a retrospective review and document outcomes. Feedback from IB practitioners, where possible, will shape subsequent cycles. The iteration cycle follows a consistent pattern: plan what to build, build it, test it, document what I learned, and decide what to do next.

I will use Git with a simplified GitFlow model (main, develop, feature branches). Continuous integration via GitHub Actions will run tests and linting on every push, enforcing 80% coverage for core modules. I will maintain a decision log using Architecture Decision Records (ADRs) to document technical choices and trade-offs. Each ADR captures the context, the decision, the alternatives considered, and the consequences. This creates an audit trail that explains not just what I built but why I built it that way.

## Requirements Analysis

Good requirements are needed to measure whether the project succeeded. This section separates business requirements (the value delivered) from functional requirements (what the system does) and non-functional requirements (how well it performs).

I gathered requirements from several sources. The problem analysis in the introduction identified core pain points from industry reports and academic literature. My own experience working in financial services informed the technical constraints. Published research on document generation and agentic AI systems shaped the functional approach. The requirements below synthesise these inputs into testable statements.

Requirements follow the MoSCoW prioritisation scheme. "Must have" requirements define the minimum viable product. "Should have" requirements add important functionality that time permitting will be included. "Could have" requirements represent stretch goals. "Won't have" requirements clarify scope boundaries.

## Business Requirements

Business requirements describe what value the project should deliver:

1. **BR1: Productivity Enhancement:** Reduce time spent on initial PB drafting by generating solid first drafts from minimal input.
2. **BR2: Knowledge Codification:** Capture institutional presentation standards through template analysis, reducing reliance on tacit knowledge.
3. **BR3: Quality Consistency:** Follow corporate visual identity standards consistently, reducing revision cycles caused by formatting issues.

4. **BR4: Data Accuracy:** Financial data in generated PBs should accurately reflect source information.

## Functional Requirements

Table 2 lists functional requirements by component using MoSCoW prioritisation (Must have, Should have, Could have, Won't have).

ID	Component	Requirement	Priority
FR1	Template Analyser	Extract slide layouts from reference .pptx files	Must
FR2	Template Analyser	Identify colour palettes and font specifications	Must
FR3	Template Analyser	Map placeholder positions and content types	Must
FR4	Data Retrieval	Fetch company financials from Yahoo Finance API	Must
FR5	Data Retrieval	Retrieve SEC filings via EDGAR API	Should
FR6	Data Retrieval	Perform web searches for company news	Should
FR7	Content Planner	Generate slide-by-slide content specifications	Must
FR8	Content Planner	Adapt content structure to PB type	Must
FR9	Content Planner	Maintain narrative coherence across slides	Should
FR10	Slide Builder	Assemble slides matching template layouts	Must
FR11	Slide Builder	Populate placeholders with generated content	Must
FR12	Slide Builder	Generate charts from financial data	Could
FR13	Orchestration	Coordinate multi-step generation workflow	Must
FR14	Orchestration	Handle API failures gracefully	Should
FR15	RAG Search	Query precedent material repository	Could

Table 2: Functional Requirements Specification

## Non-Functional Requirements

Table 3 lists quality requirements with measurable acceptance criteria.

ID	Category	Requirement	Acceptance Criterion
NFR1	Performance	System generates complete PB draft	Within 5 minutes of input submission
NFR2	Performance	API response handling	Timeout after 30 seconds with graceful degradation
NFR3	Reliability	System availability during demonstration	95% uptime during evaluation period
NFR4	Usability	Input specification interface	Requires no technical expertise to operate
NFR5	Maintainability	Codebase documentation	All public functions documented with docstrings
NFR6	Security	API credential management	No credentials in source code; environment variables used
NFR7	Security	Input validation	All user inputs sanitised before processing
NFR8	Compatibility	Output format	Valid .pptx files openable in Microsoft PowerPoint

Table 3: Non-Functional Requirements Specification

## Technology Stack

I chose technologies based on four factors:

- **Task fit:** Does the tool solve the specific problem well?
- **Familiarity:** Do I already know how to use it?
- **Maturity:** Is it well-documented with active community support?
- **Maintainability:** Will it be easy to debug and extend?

## Programming Language and Framework

I chose Python as the main language because it dominates AI and ML development and has a large library ecosystem. I am also comfortable with Python from my work experience, which means I can focus on the problem rather than learning a new language. Table 4 compares Python with alternatives.

Criteria	Python	JavaScript/Node.js	Java
Benefits	Best AI/ML libraries, mature PowerPoint library ( <code>python-pptx</code> ), rapid development	Native async, same language frontend and backend	Strong typing, enterprise support
Risks	Slower execution than compiled languages, GIL limits true parallelism	AI/ML ecosystem less mature, limited PowerPoint options	Verbose code, slower development
AI/ML support	Excellent (PyTorch, LangChain, OpenAI SDK)	Growing but limited	Moderate (DJL, Langchain4j)
PowerPoint manipulation	<code>python-pptx</code> (mature, well-documented)	Limited options	Apache POI (verbose API)

Table 4: Programming Language Comparison

FastAPI will handle the backend. It processes requests asynchronously, which matters when making multiple API calls to external data sources at the same time. When the system needs company financials and SEC filings simultaneously, async processing lets both requests run in parallel rather than waiting for one to complete before starting the other. FastAPI also generates OpenAPI documentation automatically, which will make testing easier and provides a clear contract for how the API behaves.

## AI and LLM Integration

The system will use LLMs for content planning and generation. Table 5 compares the options I considered.

Criteria	GPT-4/GPT-4o	Claude 3.5	Gemini Pro
Benefits	Best structured output, mature agent SDK, extensive documentation	Largest context window (200K), strong reasoning	Massive context (1M tokens), competitive pricing
Risks	Higher cost per token, rate limits on free tier	Newer agent tooling, less community examples	Less mature structured output, fewer agent examples
Structured output	Excellent (native JSON mode)	Good	Good
Function calling	Native, well-documented	Native support	Native support
Agent frameworks	OpenAI Agents SDK (mature)	Claude Agent SDK	Primarily via LangChain

Table 5: Large Language Model Comparison

I chose OpenAI's GPT-4o as the primary LLM. It handles structured output well and has good agent framework support through the OpenAI Agents SDK, which provides tool-use patterns for the agentic data retrieval layer.

The choice of LLM affects both capability and cost. GPT-4o represents a balance. It is capable enough to generate coherent financial content and handle multi-step reasoning for content planning. It costs more than smaller models but less than the most capable options. For a POC operating within a student budget, this balance matters. I may experiment with other models during development to understand the trade-offs, but GPT-4o will be the baseline for evaluation.

## Data Sources and APIs

I will use publicly accessible APIs to avoid proprietary data licensing complications. Proprietary data sources like Bloomberg or Refinitiv would provide richer data but come with licensing restrictions that complicate academic work. Public APIs provide enough data to demonstrate the concept while keeping the project reproducible by others. Table 6 summarises the data sources.

Source	Benefits	Risks	Access
Yahoo Finance	Free, comprehensive financial data, easy Python library (yfinance)	Unofficial API may change without notice, data delays	yfinance library
SEC EDGAR	Official regulatory filings, reliable, free	US companies only, complex document parsing required	REST API (10 req/s)
Web Search	Current news, market commentary	Cost per query, relevance filtering needed	SerpAPI or similar

Table 6: External Data Sources

## Frontend and Deployment

Table 7 compares frontend framework options.

Criteria	Next.js	React + TypeScript	Vue.js	Svelte
Benefits	Built-in API routes, server-side rendering, industry standard, strong community	Large ecosystem, strong typing, excellent tooling	Gentle learning curve, good documentation	Smaller bundle size, less boilerplate
Risks	Opinionated structure, Vercel-centric tooling	Boilerplate overhead, frequent ecosystem changes	Smaller job market, fewer enterprise examples	Immature ecosystem, limited component libraries
Component libraries	All React libraries (MUI, Chakra, Radix)	Extensive (MUI, Chakra, Radix)	Good (Vuetify, Quasar)	Limited options
TypeScript support	Native, excellent	Native, excellent	Good	Good
Deployment options	Vercel (optimised), Netlify, AWS	Vercel, Netlify, AWS	Similar options	Similar options

Table 7: Frontend Framework Comparison

I chose Next.js because it has become the industry standard for React applications. It includes built-in API routes, which simplifies the architecture by letting me handle simple backend endpoints without a

separate server. The framework has strong community support and extensive documentation. TypeScript integration is native, and I can use the full React component library ecosystem.

Table 8 compares deployment and hosting options.

Criteria	Vercel + Render	AWS	Self-hosted
Benefits	Zero configuration, automatic deployments, free tiers suitable for POC	Full control, scalable	Complete control, no vendor lock-in
Risks	Vendor lock-in, cold starts on free tier	Complex setup, cost management overhead	Time-consuming setup, maintenance burden
Cost for POC	Free tier sufficient	May incur charges	Server costs
Setup time	Minutes	Hours to days	Days

Table 8: Deployment Platform Comparison

The frontend will deploy through Vercel. The Python backend will run on Render with managed hosting. Supabase will provide PostgreSQL database services. These choices prioritise development speed over infrastructure control, which makes sense for a POC.

## Reproducibility

Reproducibility means someone else can take my work, run it in their own environment, and get the same results. This matters for two reasons. First, it lets others verify my findings rather than taking them on trust. Second, it means I can return to earlier experiments and understand exactly what produced a particular output.

LLM-based systems make reproducibility harder than traditional software. The same prompt can produce different outputs on different runs, and model providers update their systems without notice. I cannot guarantee identical results, but I can make it possible to get close.

To support this, I will pin all Python dependencies to exact versions and provide a Dockerfile so others can replicate my environment. For evaluation runs, I will set LLM temperature to zero to reduce randomness. I will separate configuration from code using a config file, and log all settings at startup. Each evaluation run will be tied to a specific Git commit, and I will log the full prompts and responses for LLM calls so I can trace how outputs were generated.

## Evaluation Methodology

I will assess the system against four metrics:

- **Efficiency:** target under 300 seconds for a 10-slide deck

- **Template fidelity:** programmatic comparison of layout, colour, and font properties
- **Data accuracy:** verification against source APIs
- **Content quality:** rubric-based assessment of narrative coherence and IB conventions

To contextualise performance, I will compare outputs against a naive baseline using generic tools (Gamma, Beautiful.ai) and estimated manual creation time. This comparison shows whether the system actually improves on existing approaches.

Ablation studies will test each component's contribution. I will run the system with different components disabled: content generation without template analysis, template analysis without LLM planning, and data retrieval without web search. These experiments will reveal which parts of the pipeline contribute most to output quality.

Evaluation will use fixed test cases covering different company types, sectors, and PB types. The test set will include large-cap technology companies, mid-cap manufacturing firms, and small-cap financial services companies. Each test case will be documented with input parameters and expected outputs to support replication.

For qualitative assessment, I will develop a rubric covering content accuracy, narrative coherence, professional tone, and formatting correctness. If I can access IB practitioners for feedback, their input will supplement the rubric-based evaluation. However, the core evaluation does not depend on external participation.

## Project Plan

The project will run for thirteen weeks from start to final submission. I have divided it into phases matching the iterative approach. The schedule balances implementation time against documentation requirements, recognising that academic projects require substantial written output alongside working software.

The first phase focuses on research and planning. The literature review and methodology chapters establish the theoretical foundation. This front-loaded documentation ensures that implementation decisions rest on solid ground.

The middle phase covers implementation in increasing depth. I will start with the foundational components (template analysis, data retrieval) before tackling integration and orchestration. This sequencing means I will have working pieces to test before attempting to connect them.

The final phase covers evaluation and finalisation. I have allocated time for proper analysis rather than rushing to document results at the last minute.

Figure 1 shows the schedule as a Gantt chart.

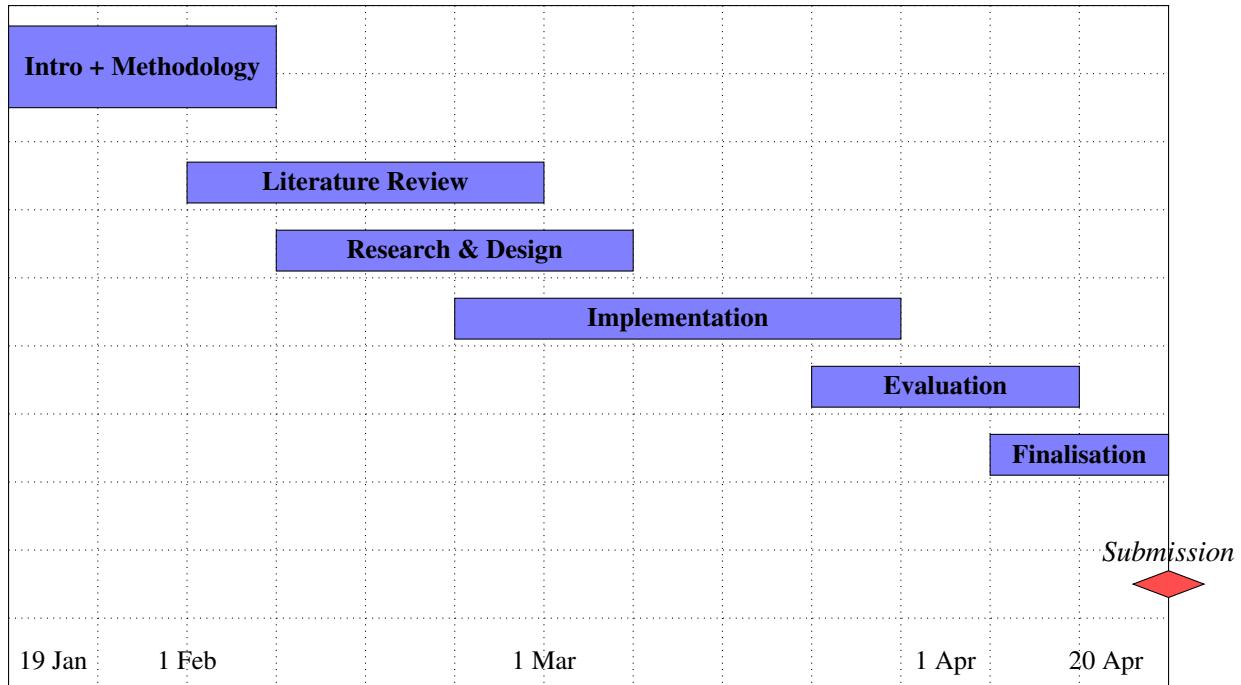


Figure 1: Project Gantt Chart (19 January to 20 April 2026)

Table 9 lists the milestones with their specific deliverables and target dates. Each milestone represents a checkpoint where I will assess progress and decide whether to continue as planned or adjust course. The dates are targets rather than commitments. If a milestone slips, I will evaluate whether to compress subsequent work, reduce scope, or accept schedule delay.

Phase	Date	Milestone	Deliverables
1	19 Jan	Project Initiation	Repository setup, development environment configured
2	31 Jan	Initial Chapters	Introduction, Scope, Methodology chapters complete
3	21 Feb	Literature Review	Literature review chapter complete
4	28 Feb	Design Complete	Architecture diagrams, API specifications, data models
5	14 Mar	Core Components	Template analyser and data retrieval functional
6	28 Mar	MVP Complete	End-to-end generation pipeline operational
7	4 Apr	Testing Complete	Unit and integration tests passing; evaluation data collected
8	11 Apr	Evaluation Complete	Results analysis and evaluation chapter written
9	18 Apr	Document Finalised	All chapters complete, proofread, formatted
10	20 Apr	Submission	Final report submitted via QMPlus

Table 9: Project Milestones and Deliverables

## Risk Assessment and Mitigation

Every project faces risks. The question is whether to identify them upfront and plan responses, or to discover them later when options are limited. I prefer the former.

I have categorised risks by source: technical risks (API limitations, LLM behaviour, template complexity), process risks (scope creep, time pressure), and external risks (API changes, compliance requirements). The mitigations focus on practical actions I can take rather than hoping problems do not occur.

Table 10 shows the risk register with likelihood and impact rated on a three-point scale (Low, Medium, High).

ID	Risk Description	Likelihood	Impact	Mitigation Strategy
R1	API rate limits restrict data retrieval during development	Medium	Medium	Implement caching; use mock data for testing; stagger API calls
R2	LLM outputs inconsistent or hallucinated content	High	High	Structured output schemas; validation layers; HITL review
R3	Template analysis fails on complex slide layouts	Medium	High	Scope to common layouts; graceful degradation for unsupported elements
R4	External API deprecation or changes	Low	High	Abstract API interactions; monitor changelogs; maintain fallback sources
R5	Scope creep extends beyond POC boundaries	Medium	Medium	Strict MoSCoW prioritisation; regular scope reviews against requirements
R6	Technical complexity exceeds available time	Medium	High	Iterative delivery; prioritise core pipeline; defer stretch goals
R7	Data accuracy issues undermine credibility	Medium	High	Verification against manual retrieval; source attribution in outputs
R8	Generated outputs violate compliance requirements	Low	High	Human review mandatory; disclaimer on all outputs; no PII processing

Table 10: Risk Assessment and Mitigation Strategies

For high-impact risks, I have simple contingency rules. If LLM validation failures become frequent, I will reduce scope. If template extraction struggles with complex layouts, I will narrow the supported template types. If milestones slip by more than a week, I will drop the "Could have" requirements first.

## Ethics, Data Governance and Compliance

AI systems in financial services raise ethical concerns worth addressing directly. The documents they help produce can influence major business decisions. Errors or bias in generated content could mislead clients or expose firms to liability. The design must address these concerns directly.

The system keeps humans in control of analytical judgements. Generated content is marked as AI-assisted, and HITL design ensures professional review before distribution. The system prioritises accuracy over completeness, flagging uncertainty rather than presenting unverified information as fact. When the system cannot retrieve data or has low confidence in generated content, it says so explicitly rather than filling gaps with plausible-sounding fabrications.

For data governance, the system processes only data needed for PB generation, with no retention of user

inputs after sessions. All external data includes source attribution. The POC does not process personally identifiable information. Company data comes exclusively from public sources. This approach sidesteps many privacy concerns while still demonstrating the core functionality.

Regarding compliance, the HITL architecture meets EU AI Act requirements for human oversight (Bank for International Settlements 2024). Professional responsibility for final content remains with the banker. The system logs generation parameters and data sources to support audit requirements. If questions arise about how a particular output was generated, the logs provide a complete trail from input to output.

## **Project Tracking**

Progress will be tracked using GitHub Issues with a Kanban board (Backlog, In Progress, In Review, Done). Each functional requirement maps to one or more issues, providing traceability from requirements to implementation. Changes to scope or architecture will be documented in the issue history.

This approach serves two purposes. First, it keeps me organised during development. I can see at a glance what is done, what is in progress, and what remains. Second, it creates an audit trail for the dissertation. Anyone reviewing the work can trace how implementation decisions connected to requirements and how the project evolved over time.

## **Reusability**

Academic projects should produce reusable artefacts where possible. Code that only works for one specific case has limited value.

The architecture separates concerns into independent modules (template analysis, data retrieval, content planning, slide assembly), each usable on its own. Someone interested only in template analysis could use that module without the rest of the system. Someone building a different kind of document generator could reuse the data retrieval layer.

All code will include Google-style docstrings, with a detailed README and architecture documentation. The codebase will be released under MIT License to minimise barriers to reuse. Configuration will be externalised so that adapting the system to different use cases does not require modifying source code.

# Preliminary Research and Design Documentation

Source	Type	Summary	Link
<b>Knowledge Management &amp; Knowledge Reuse</b>			
Markus (2001). Toward a Theory of Knowledge Reuse	Journal	Foundational theory on knowledge reuse; identifies four types of reuse situations and success factors	<a href="#">Link</a>
Dalkir (2017). Knowledge Management in Theory and Practice	Textbook	Comprehensive KM textbook covering tacit/explicit knowledge and organisational memory systems	<a href="#">Link</a>
<b>Large Language Models</b>			
Zhao et al. (2023). A Survey of Large Language Models	Survey	Comprehensive survey on LLM development; covers GPT, LLaMA, PaLM families	<a href="#">Link</a>
Minaee et al. (2024). Large Language Models: A Survey	Survey	Reviews LLM characteristics, contributions, limitations, and augmentation techniques	<a href="#">Link</a>
<b>Retrieval-Augmented Generation</b>			
Lewis et al. (2020). RAG for Knowledge-Intensive NLP Tasks	Conference	Seminal RAG paper; combines parametric and non-parametric memory for factual generation	<a href="#">Link</a>
Gao et al. (2023). RAG for LLMs: A Survey	Survey	Comprehensive RAG survey covering taxonomy, methods, and applications	<a href="#">Link</a>
Chen et al. (2025). RAG and LLMs for Enterprise KM	Journal	Systematic review of 63 studies; 63.6% use GPT models, 80.5% use FAISS/Elasticsearch	<a href="#">Link</a>
Fan et al. (2024). A Survey on RAG Meeting LLMs	Conference	KDD survey on integrating RAG with LLMs	<a href="#">Link</a>
<b>Agentic AI &amp; Autonomous Agents</b>			
Wang et al. (2024). A Survey on LLM-based Autonomous Agents	Journal	Foundational survey on LLM autonomous agents; proposes unified framework	<a href="#">Link</a>
Guo et al. (2024). LLM Based Multi-agents: A Survey	Conference	IJCAI survey on multi-agent LLM systems for complex problem-solving	<a href="#">Link</a>
Li et al. (2024). A survey on LLM-based multi-agent systems	Journal	Systematic review with five-component architecture framework	<a href="#">Link</a>
<b>LLM Tool Use &amp; Function Calling</b>			
Li (2025). A review of paradigms for LLM-based agents	Conference	Reviews tool use, planning, RAG, and feedback mechanisms	<a href="#">Link</a>
<b>Automated Document Generation</b>			
Zheng et al. (2025). PPTAgent: Generating Presentations	Conference	Two-stage edit-based presentation generation; introduces PPTEval framework	<a href="#">Link</a>

Source	Type	Summary	Link
Liu et al. (2024). We Need Structured Output	Conference	Studies user needs for structured LLM outputs in professional contexts	<a href="#">Link</a>
<b>LLM Structured Output &amp; Prompt Engineering</b>			
Wu et al. (2024). LLM-Driven Structured Output: A Benchmark	Journal	Benchmark for structured outputs; compares fine-tuning, prompting, and RAG	<a href="#">Link</a>
Xu et al. (2025). Structured Data Generation with GPT-4o	Journal	Compares JSON, YAML, CSV prompt styles; JSON best for complex data	<a href="#">Link</a>
<b>LLM Hallucination &amp; Factual Accuracy</b>			
Huang et al. (2024). A Survey on Hallucination in LLMs	Journal	Comprehensive taxonomy of hallucination causes, detection, and benchmarks	<a href="#">Link</a>
Alansari & Luqman (2025). LLM Hallucination Survey	Survey	Review of hallucination causes, detection, and mitigation strategies	<a href="#">Link</a>
<b>Human-in-the-Loop AI</b>			
Wu et al. (2022). A Survey of HITL for ML	Journal	Classifies HITL approaches: data processing, interventional training, system design	<a href="#">Link</a>
Mosqueira-Rey et al. (2024). HITL ML: Role of the user	Journal	Discusses timing, frequency, and workload factors in interactive ML	<a href="#">Link</a>
Rezaeighaleh et al. (2023). Ethical AI Based on HITL	Journal	Examines HITL for ethical AI development	<a href="#">Link</a>
<b>AI in Financial Services</b>			
Alghofaili et al. (2024). AI and ML in Banking Systems	Journal	Examines board role in AI adoption; Saudi Arabian banking sector	<a href="#">Link</a>
Kumari & Tanwar (2023). AI/ML in Financial Services	Journal	Bibliometric analysis of 1,045 BFSI sector articles	<a href="#">Link</a>
<b>AI Regulation &amp; Compliance</b>			
Zetsche et al. (2024). Regulating AI in investment management	Journal	Discusses EU AI Act, GDPR, MiFID II implications for AI	<a href="#">Link</a>
<b>Knowledge Worker Productivity</b>			
Brynjolfsson et al. (2023). Generative AI at Work	Working Paper	14% productivity increase from AI; greatest impact on novice workers	<a href="#">Link</a>
Noy & Zhang (2023). Productivity effects of generative AI	Journal	Experimental study showing productivity improvements from ChatGPT use	<a href="#">Link</a>
McKinsey (2012). The Social Economy	Report	Knowledge workers spend 20% of time searching for information	<a href="#">Link</a>
<b>Vision-Language Models for Documents</b>			

Source	Type	Summary	Link
Faysse et al. (2024). ColPali: Document Retrieval with VLMs	Research	Uses VLM for direct PDF image retrieval; introduces ViDoRe benchmark	<a href="#">Link</a>
Liao et al. (2024). DocVLM: Efficient Reader	Conference	Integrates OCR into VLMs; improves DocVQA from 56% to 86.6%	<a href="#">Link</a>
<b>Software Engineering Methodology</b>			
Anifa et al. (2024). Systematic Review on Agile Approach	Journal	Systematic review of agile methodology across industries	<a href="#">Link</a>
Sommerville (2016). Software Engineering	Textbook	Standard software engineering reference; covers SDLC and methodologies	<a href="#">Link</a>

## **Project Implementation and Outcomes**

## **Evaluation and Conclusions**

# References

- Bank for International Settlements (2024). *Regulating AI in the Financial Sector: Recent Developments and Main Challenges*. 63. URL: <https://www.bis.org/fsi/publ/insights63.pdf>.
- BBC News (Mar. 2021). *Goldman Sachs junior bankers rebel over ‘inhumane’ 100-hour weeks*. URL: <https://www.bbc.co.uk/news/business-56452494> (visited on 01/15/2025).
- Corporate Finance Institute (2025). *Investment Banking Analyst Job Description, Hours, and Salary*. URL: <https://corporatefinanceinstitute.com/resources/career/investment-banking-analyst-job-description-hours-salary/> (visited on 02/24/2025).
- Dalkir, Kimiz (2011). *Knowledge Management in Theory and Practice*. 2nd. Cambridge, MA: MIT Press. ISBN: 978-0262015080.
- IBM (2024). *What are Agentic Workflows?* URL: <https://www.ibm.com/think/topics/agentic-workflows> (visited on 11/18/2025).
- Markus, M. Lynne (2001). “Toward a Theory of Knowledge Reuse: Types of Knowledge Reuse Situations and Factors in Reuse Success”. In: *Journal of Management Information Systems* 18.1, pp. 57–93. DOI: 10.1080/07421222.2001.11045671.
- McKinsey & Company (2024). *Capturing the Full Value of Generative AI in Banking*. McKinsey & Company. URL: <https://www.mckinsey.com/industries/financial-services/our-insights/capturing-the-full-value-of-generative-ai-in-banking>.
- McKinsey Global Institute (July 2012). *The Social Economy: Unlocking Value and Productivity Through Social Technologies*. McKinsey & Company. URL: <https://www.mckinsey.com/industries/technology-media-and-telecommunications/our-insights/the-social-economy>.
- Radhakrishna, Vedapradha et al. (2024). “Future of Knowledge Management in Investment Banking: Role of Personal Intelligent Assistants”. In: *SAGE Open* 14.4. DOI: 10.1177/20597991241287118.
- Sommerville, Ian (2016). *Software Engineering*. 10th. Boston, MA: Pearson Education. ISBN: 978-0133943030.
- Wall Street Oasis (2024). *Investment Banking Hours: Typical Schedule, Work-Life Balance, and Expectations*. URL: <https://www.wallstreetoasis.com/resources/careers/jobs/investment-banking-hours>.
- Wang, Lei et al. (2024). “A Survey on Large Language Model based Autonomous Agents”. In: *Frontiers of Computer Science* 18.6, p. 186345. DOI: 10.1007/s11704-024-40231-1. arXiv: 2308.11432.
- Zheng, Hao et al. (2025). “PPTAgent: Generating and Evaluating Presentations Beyond Text-to-Slides”. In: *arXiv preprint*. arXiv: 2501.03936. URL: <https://arxiv.org/abs/2501.03936>.

## Appendices